


# Parallel Computing Through Code Analysis

Clark Fitzgerald

UC Davis

10 May 2017



# Table of Contents

- 1 Introduction
- 2 Motivating Example
- 3 Parallel Concepts
- 4 Code Analysis
- 5 Conclusion

Modern platforms provide incredible computing power.




But they require expertise.

The broader goal is for users to write higher level code that also performs better.

- Parallel programming is a means to this end
- Compilation is another way
- Expertise in system rather than end user

We take a holistic approach to the computation.



The R language offers several benefits.



- Functional languages simplify parallel computing
- Widely used for statistics and data analysis
- Supports metaprogramming aka “programming on the language”

# Table of Contents


- 1 Introduction
- 2 **Motivating Example**
- 3 Parallel Concepts
- 4 Code Analysis
- 5 Conclusion

The purpose of this example is to motivate the proposed research.

- Working with Professor Michael Zhang from Civil Engineering
- Illustrates complexity when computing with larger data sets



Loop detectors count vehicle flow, measuring velocity and density (time sensor is activated).




# Caltrans Performance Measurement System (PeMS)


records loop detector data for the whole state.

- Each sensor measures 3 quantities
- Data point every 30 seconds
- 43,680 sensors in California
- $\implies$  377 million data points per day

The *fundamental diagram* in traffic engineering shows the relationship between flow and density.




Each station has a fundamental diagram, which can be fit in parallel.



Source: Li, Zhang 2011

Using more data allows new types of analysis.

## R expresses statistical computation well.



```
by(data=single_day, INDICES = station, FUN = my_fd)
```

- For a single day with 377 million observations this can be done on a single machine.
- A sensible way to run this in parallel is to `fork()` the process after reading in the data.
- So you write a bunch of code to do that :)

## Small changes can require totally different computations.

If we compute on one year then this will exceed memory.

```
by(data=one_year, INDICES = station, FUN = my_fd)
```

A different model, such as least squares, may be able to process the data as a stream.

```
by(data=one_year, INDICES = station, FUN = my_lm_fd)
```

Access to the underlying database may allow us to run code directly inside the database.

```
SELECT station, my_fd(...) FROM data GROUP BY station
```

# Table of Contents

- 1 Introduction
- 2 Motivating Example
- 3 Parallel Concepts**
- 4 Code Analysis
- 5 Conclusion



# This simple example shows how to write parallel code in R.

Consider computing the mean,

$$\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i \quad (1)$$

where the  $x_i$ 's are i.i.d.  $\sim t(d)$ .

In R this code is written:

```
xbar = mean(rt(n, d))
```

We can express the mean as a weighted mean.

Suppose  $n = n_j p$ , where  $n_j$  is the chunk size and  $p$  is the number of chunks.


$$\bar{x} = \frac{1}{n} \sum_{j=1}^p \sum_{i=1}^{n_j} x_{ij} = \frac{1}{p} \sum_{j=1}^p \frac{1}{n_j} \sum_{i=1}^{n_j} x_{ij} = \frac{1}{p} \sum_{j=1}^p \bar{x}_{\cdot j} \quad (2)$$

The weighted mean can be directly translated into R code.

```
partial_means = replicate(p, mean(rt(n_j, d)))  
xbar = mean(partial_means)
```

- While not parallel, this effectively removes the memory limits.
- How to choose  $n_j$  and  $p$ ?

The same computation can be evaluated on many workers simultaneously.



Here is one way to parallelize this code.

```
library(parallel)
p = floor(detectCores(logical = FALSE) / 2)
n_j = n / p
cluster = makeCluster(p)

expr = substitute(mean(rt(n_j, d)),
list(d = d, n_j = n_j))

partial_means = unlist(
  clusterCall(cluster, eval, expr))

xbar = mean(partial_means)
```

We're considering a system that transforms expressions.

**Input:**

```
xbar = mean(rt(n, d))
```

**Output:** (omitting boilerplate)

```
p = floor(detectCores(logical = FALSE) / 2)
```

```
expr = substitute(mean(rt(n_j, d)),  
  list(d = d, n_j = n_j))
```

```
partial_means = clusterCall(cluster, eval, expr)  
xbar = mean(partial_means)
```

This can be difficult because R is implemented in C.

```
> head(rt, 4)
```

```
1 function (n, df, ncp)
2 {
3     if (missing(ncp))
4         .Call(C_rt, n, df)
```

Options:

- Start from `replicate(p, mean(rt(n_j, d)))`
- Allow users to indicate how `rt` is vectorized
- Analyze the preprocessed C code
- Rewrite the C code in R, then analyze the R code

# Table of Contents

- 1 Introduction
- 2 Motivating Example
- 3 Parallel Concepts
- 4 Code Analysis**
- 5 Conclusion

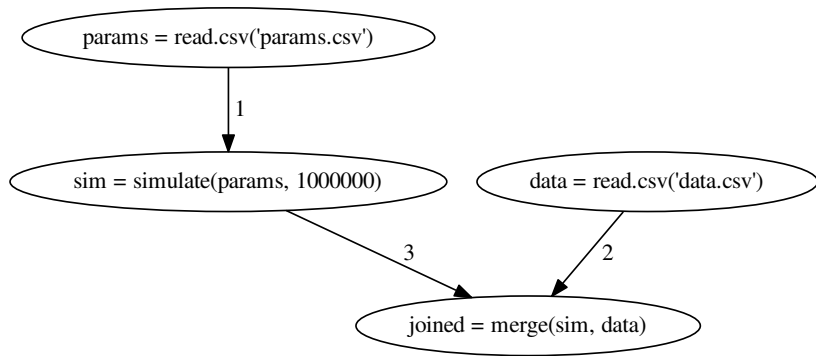


CodeDepends is a tool for analyzing code as a data structure.

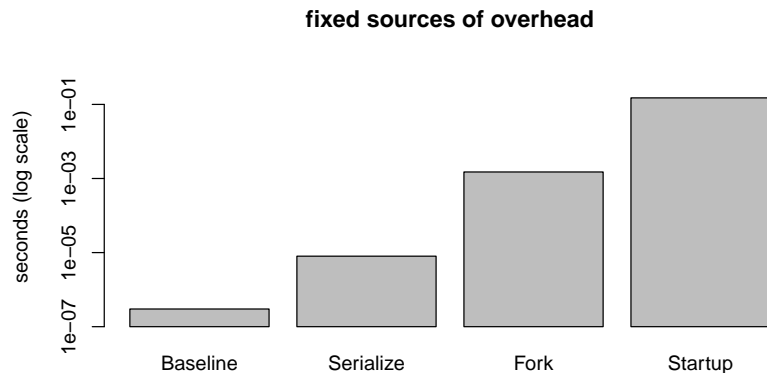
Consider this script:

```
params = read.csv('params.csv')
data = read.csv('data.csv')
sim = simulate(params, 1000000)
joined = merge(data, sim)
```

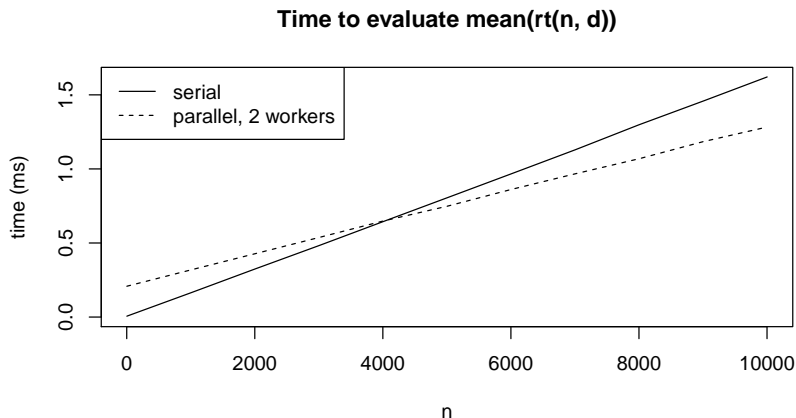
The expression graph represents the dependencies between expressions.



# Is it worth it to go parallel?



Given an existing SNOW cluster with 2 workers we see benefits from parallelization when  $n > 4000$ .



Timings on a 3.4 GHz Intel i3 CPU

# Factors to consider

## Parameters

- number of processor cores to use
- size of each chunk
- which functions to combine in one processing step

## Constraints

- number of cores available
- network bandwidth
- disk IO speed
- available memory

Idiomatic R already expresses computation in a natural parallel way through “apply” functions.

```
x = replicate(5, rnorm(n_j), simplify = FALSE)
```

```
partialmeans = lapply(x, mean)
```

```
by(data, INDICES = station, FUN = piecewise_rlm)
```


Also

```
apply, sapply, tapply, by, mapply, Map, vapply, outer
```

Layers mark ways for users to write parallel code for one platform.

- **User R Packages:** foreach, future, partools, ddR, biganalytics, RevoScaleR
- **R Layer:** SNOW, parallel, bigmemory, Rmpi, RCUDA
- **Operating System:** threads, processes, \*NIX fork(), memory maps, network sockets, MPI

# How can we transform R code into a lower layer?





Knowledge of the data allows us to generate more specialized code.

- File size
- Dimensions of table / matrix / array
- Column classes
- Randomized rows
- Sorted / grouped
- Possible values for factor
- Indexed
- Including sufficient statistics

## Example: a data format that facilitates sampling.

station, flow, occupancy, time

1	12	0.087	09:57:00
---	----	-------	----------

1	14	0.092	14:29:30
---	----	-------	----------

...

7	14	0.088	16:32:30
---	----	-------	----------

7	11	0.090	17:12:00
---	----	-------	----------

- ASCII fixed width format,  $c$  characters (bytes) per row
- sorted on station, then occupancy
- $r$  rows per station
- $\implies$  new stations begin at byte  $i \times c \times r$

# Table of Contents

- 1 Introduction
- 2 Motivating Example
- 3 Parallel Concepts
- 4 Code Analysis
- 5 Conclusion

We propose using {Code, Data, Platform} to determine a parallel execution strategy.

```
1 library(BenfordTests)
2
3 # Using data downloaded from
4 # http://www.fec.gov/finance/disclosure/ftpdet.shtml
5 # Individual contributions file for 2017 is 4GB
6
7 r = edist.benftest(x)
8
9 header = as.character(read.table("indiv_header_file.csv", sep = ",",
10                                stringsAsFactors = FALSE))
11
12 keepers = c("CITE_ID", "ENTITY_ID", "TRANSACTION_AMT")
13 keep = header %>% keepers
14
15 cc = rep("NULL", length(header))
16 cc[keep] = c("factor", "factor", "numeric")
17
18 ltcont = read.table("~/data/ltcont.txt",
19                    sep = "\t",
20                    colClasses = cc,
21                    nrows = 200)
22
```



## Related Work

- Related CS literature mostly focused on algorithmic applications and lower level languages
- Software such as Dask, Theano, Tensorflow require the user to build computation graphs
- Primarily focused on numeric arrays

# The next step is to build a prototype of the system.

Specifically beginning with:

- *code*: Apply family of functions
- *data*: Files on disk exceeding main memory
- *platform*: Single server

Then test it with the traffic sensor data.

Next I'd like to extend this to platforms that store data

```
SELECT station, my_rlm(...) FROM data GROUP BY station
```


- Better data locality
- Many use cases
- Connection to compilation

# Acknowledgements

- Faculty members
- Data Science Institute Affiliates and statistics students for applications and feedback
- Special thanks to Professors Duncan Temple Lang and Michael Zhang



# Questions?



# Additional Slides

## Follow up on the sensor example

The transformed program should:

- ➊ Remove any observations it can
- ➋ Reorganize files on disk based on station ID
- ➌ Apply function to each station ID file

# More Applications

- Benford test on election campaign contribution data
- Forest greenness satellite imagery (Andrew Latimer)
- Simulating spread of disease (Nistara Randhawa)

## Preserving language semantics can be challenging.

For example, R's dynamic lookups

```
f = function() 0
g = function() f() + 1
f = function() 10
g()                                # Returns 11!
```

## Compiled R code provides even more efficiency.

- ① Parallelization will complement efforts to compile R
- ② Compiled code potentially allows the use of shared memory threads
- ③ May follow the OpenCL kernel model

## Last summer I worked on the Distributed Data Structures in R (DDR) project

- Relevant experience
- Idea: an abstraction layer for distributed and parallel data structures
- Created R lists and apply type functions to run on Spark

## How do we detect if a function in R is vectorized, and in which arguments?

`rnorm()` is vectorized in the last two arguments, but semantically different for a vector in the first argument.


```
> rnorm(5, mean = c(1, 2), sd = c(2, 10, 200))  
[1] 0.2134756 -4.1137221 256.0094734 0.4562226 -10.385537
```

It's all C Code.

```
> rnorm  
function (n, mean = 0, sd = 1)  
.Call(C_rnorm, n, mean, sd)
```



# An iterator produces data on demand



- Most flexible of the above options
- Natural in pipeline parallel model
- Operate well with high performance IO libraries
- Unfamiliar to R programmers
- Not ideal if you need the whole data set

## R in a database

Suppose we want to call the vectorized function  $f()$ . If  $f()$  is available both in R and the database then we have two options:

Option 1: Run  $f$  inside database, returning result

```
dbGetQuery(con,"SELECT f(x) FROM mydata;")
```

Option 2: First fetch  $x$ , then call  $f$  within R


```
f(dbGetQuery(con,"SELECT x FROM mydata;"))
```

By “programming on the language” we can modify existing code.

```
lapply_to_mclapply = function(expr)
{
  # Changes lapply to parallel::mclapply
  lapply = quote(parallel::mclapply)
  expr = force(expr)
  # Following Wickham's Advanced R book
  call = substitute(substitute(expr))
  eval(call)
}
```

```
> e1 = quote(xmeans <- lapply(x, mean))
> lapply_to_mclapply(e1)
xmeans <- parallel::mclapply(x, mean)
```

## Pipeline parallelism is like a factory assembly line.



```
# Worker 1
```

```
x_chunk = rnorm(n_j)
```

```
serialize(x_chunk, worker2)
```

```
# Worker 2
```

```
x_chunk = unserialize(worker1)
```

```
partial_means[i] = mean(x_chunk)
```

Simple row based sampling misses the important areas of high density.

